MINI-HOW-TO

MOCASSIN Version 2.02

(MOnte CArlo SimulationS of Ionized Nebulae)

LAST MODIFIED: 19/09/2008

barbara ercolano, harvard-smithsonian centre for astrophysics, cambridge, ma

CONTENTS

0. What's new

from 2.02.52 to 2.02.53 Updated atomic data to Chianti 6. Also added NiII and CaII from the forthcoming Chianti 7. Changed the SED.out output to give flux in Jy multiplied by D^2 where D is in pc (users must divide by D^2 to obtain flux in Jy. Reinstated the accessories directory with a plotSED.pro idl program included. Echo keyword included - see list of keywords below.

from 2.02.51 to 2.02.52 Fixed bug in mocassinWarm and mocassinOutput in grid_mod and fixed but that affected plane parallel models.

from 2.02.50 to 2.02.51 Fixed bug in mocassinOutput mocassinWarm

from 2.02.49 to 2.02.50 2D option implemented. Does not work with multigrids.

from 2.02.48 to 2.02.49 removed the tau.out output, now mocassin produces tauNu.out output, which is the tau(nu) at the edge of the grid in the three axial directions starting from the origin of the axes.

from 2.02.47 to 2.02.48 removed bad Chianti data and reformatted data comments to be more portable.

from 2.02.46 to 2.02.47 the lgIsotropic option is now carried through to the mocassinWarm start routine

from 2.02.45 to 2.02.46 added unisotropic scattering - the ¡cos theta¿ are calculated directly from the dielectric constants through MIE theory and make use of the Heyney Greenstein phase function - polarisation not included but would be trivial at this point. Isotropic scattering case is still available via the inclusion of the keyword isotropicScattering

from 2.02.44 to 2.02.45 fixed bug that caused mocassin to hang after convergence has been reached

from 2.02.43 to 2.02.44 fixed bug in HeI rec lines routine (extrapolation beyond atomic data temperature limits caused inappropriate He I recombination line ratios- only affected HeI recombiation line estimates for model with a significant number of cells with Te¡ 5kK). Bug affected versions 2.02.31-2.02.42.

from 2.02.42 to 2.02.43 fixed bug in nu searching routine (caused spikes in SED when resolution changed dramatically)

from 2.02.41 to 2.02.42 fixed bug in dust PDF calculations

from 2.02.40 to 2.02.41 fixed bug in Lphot-¿Lstar conversion for dust+gas simulations

from 2.02.39 to 2.02.40 fixed bug in HeI rec lines routine

from 2.02.38 to 2.02.39 fixed allocation bug in update - grain heating - only affected g-fortran; fixed bug affecting multistar models in symmetricXYZ run for only 1 source not centtrally located

from 2.02.37 to 2.02.38 g-fortran compatible

from 2.02.36 to 2.02.37 nForLevels can be used for FeII

from 2.02.35 to 2.02.36 bug fixed in statistical weightsoutshells affecting species with z¿18

from 2.02.34 to 2.02.35 bug fixed in reading PAH files

from 2.02.33 to 2.02.34 bug fixed in mocassinPlot

from 2.02.32 to 2.02.33 HeI recombination lines routine fixed in emission_mod and output_mod.

from 2.02.31 to 2.02.32 bugs fixed in output_mod - recombination contribution to forbidden transitions added to selected ions - ci.dat ciii.dat neiv.dat nev.dat ni.dat nii.dat niv.dat oii.dat oiii.dat (W. Wei, Hai Bo - PKU); equilibrium routine and its calls also changed accordingly; to see what CELs now include the recombination contributions please see the end of the data/*.dat files listed.

from 2.02.30 to 2.02.31 HeI recombination line data modified to include density dependancy.

from 2.02.29 to 2.02.30 forbiddenlines pointers modified.

from 2.02.28 to 2.02.29 bug fixed with mocassinPlot.

from 2.02.14 to 2.02.28 some bugs fixed, most regarding the application of multigrids and of multiple ionisation sources. Previous syntax for multigrids and multiple sources is obsolete. Please check relevant sections. New autmatic contShape allowed *powerlaw*. Plane ionisation cases: energy threshold option introduced – see planeIonization. New keywords: *nstages*, *getEquivalentTau*.

from 2.02.12 to 2.02.14 diffuse ionisation source included see diffuseSource keyword. V 2.02.13 was skipped as it includes extinction map routines - redundant - and also it has a higher number of Hlevels so it runs slower.

from 2.01 to 2.02 ionising stellar atmosphere files from http://astro.uni-tuebingen.de/ rauch can now be directly read into MOCASSIN (but please manually remove the header). Grain charge calculations. Photoelectric heating from dust grains. Gas-grain collisions. Tspiking. PAH optical constants included. Dust optical constant data can be entered in terms of Qs (lambda dependent and radius dependent). Please refer to *Input file* section of this manual.
See changes to keyword: *contShape*
See new keywords: *traceHeating*, *quantumHeatGrain*, *quantumHeatGrainParameters*, *noPhotoelectric*

from 2.00 to 2.01 The main changes regard the introduction of multiple ionising sources placed at any location in the grid. See section 3.1 *multiPhotoSources*.

The central star(s) parameters are now written to a separate output file (output/photoSource.out). This version is fully compatible with inputs of Version 2.00, but not with the grid*.out files (so mocassinWarm compiled with Version 2.01 will not work with grids produced by Version 2.00).

1. How to get MOCASSIN up and running

1.1 Obtaining the MOCASSIN source code

We have decided that MOCASSIN's source code will now only be available on request by sending an email to Barbara Ercolano (be@star.ucl.ac.uk). The reason for this is that we would like to have an idea of who is using the code, making it easier to provide info and new versions and, most importantly, should bugs emerge later on, also provide patches. If you have previously downloaded Version 1.01 or earlier, please email me to obtain the new version.

1.2. Unpack the tarball

gunzip mocassin.X.XX.tar.gzip

tar xvf mocassin.X.XX.tar

This will create the subdirectory mocassin.X.XX and all the files will be put there.
The following subdirectories will also be created:

mocassin.X.XX/source → source code for all the modules
mocassin.X.XX/input → where mocassin looks for the input files for the current simulation
mocassin.X.XX/output → where mocassin writes the output files
mocassin.X.XX/benchmarks → input and output files from the benchmark tests
mocassin.X.XX/data → containing all atomic data files needed by mocassin
mocassin.X.XX/dustData → optical constants library and other dust data files
mocassin.X.XX/examples → some examples input files
mocassin.X.XX/accessories→ miscellaneous FORTRAN and IDL programs

1.3. Compile mocassin, mocassinWarm, mocassinOutput or mocassinPlot

In the tar ball you should find an example makefile; this was coded for use with the current free Intel F90 compiler (ifort 8.0.034) and the standard free MPI libraries (mpich-1.2.5.2) Please edit the makefile to suit the path of the compiler and mpi libraries on your system. some compilers may require different/additional optimisation flags.

Several other makefiles created for a number of different platforms are also included (e.g. makefile.sun etc..); these are only given as further examples and are not being mantained.

Whatever you do PLEASE be careful with aggressive optimization switches. -O5 seems to work fine on all platforms so far, but -fsimple=2 gives serious problems. **ALWAYS RUN A BENCHMARK CASE WHEN COMPILING ON A NEW PLATFORM!!**.

There are four drivers for MOCASSIN:

mocassin: the main mocassin driver that is used to start a new simulation.

mocassinWarm: the driver that is used to resume an interrupted simulation.

mocassinOutput: runs the output routines using the current grid files in the output subdiretory.

mocassinPlot: uses the current grid files in the output/ subdirectory to create 3d-emission maps.

To compile either of these drivers you need to do type:

make mocassin
make mocassinWarm
make mocassinOutput
make mocassinPlot

WARNING!!!! Sometimes (e.g. on the SUN) the makefile is such that you have to type

make

first, this will make all the object files, then to link up one of the versions of the driver listed above type

make mocassin
make mocassinWarm
make mocassinOutput
make mocassinPlot

.... if no compiler errors and everything has gone smoothly then the sun will shine and all will be very nice.. if you get compilation errors please email them me them (bestar.ucl.ac.uk) and I will try to help if I can.

IMORTANT NOTE: If you are using an IBM platform and the makefile provided here then you will also need to change all the module file extensions from .f90 to .f (this is because for some annoying reason the compiler only looks for .f files.. I am sure there must be a way around this, but I couldn't find it, if you do, please let me know!). The file extensions can be easily changed using the following single line awk command:

ls *.f90 — awk 'gsub("90","")print"mv "1"90"1' — sh

1.4. Writing a new Makefile
MOCASSIN's source code is identical for all platforms, only the makefile changes. To write a new makefile a good place to start is to find out what Fortran 90 compiler is installed on your machine and then read the compiler's man pages. e.g. if you are using xlf90, then just type man xlf90 and read all the available compiler options. You will need to include some sort of optimisation flag and some other MPI related flags. You can refer to the makefiles given here for the other platforms.

2. How to run the benchmark cases

2.1 Pure photoionisation benchmarks

2.1.1 The benchmark problems
Once you have downloaded, unpacked and compiled MOCASSIN successfully you can first of all try to run one or more of the available benchmark problems. These were designed at a series of workshops on photoionised plasma codes held at Meudon (France) and Lexington (USA). The latest version of these benchmarks are included in Pequignot et al. (2001) in Spectroscopic Challenges of Photoionised Plasmas, ASP Conference Series Vol 247; G. Ferland and D.W. Savina eds. Should these conference proceedings not be available to you, the same benchmarks are also reported in Ercolano et al., 2003, MNRAS, 340, 1136. Here you will also find some more general info about MOCASSIN. (There is also a link to this paper at www.star.ucl.ac.uk/ be, following the refereed journal papers link). Please note that the MOCASSIN results included in the Pequignot et al. (2001) were only from a very early version of the code. Best results are those listed in the Ercolano et al., 2003, MNRAS, 340, 1136 paper.

2.1.2 Input files for benchmarks
Copy the input file of the benchmark you wish to run (you should have found these in a subdirectory called benchmarks, also included in the tar ball) into the mocassin.X.XX/input subdirectory and re-name it input.in or link the input.in file to the benchmark file.

e.g. for the Meudon standard HII region (T_{star}=40kK) (assume MOCASSIN is in /mocassin.X.XX/) type
    cp /mocassin.X.XX/benchmarks/gas/HII40/meudonHII40.in /mocassin.X.XX/input/input.in
or
    ln -s /mocassin.X.XX/benchmarks/gas/HII40/meudonHII40.in /mocassin.X.XX/input.in

also copy the respective nebular abundance file into the same directory (no need to rename!). In the case of the Meudon standard HII region benchmark this will be abunHII40.in (you can check the name of the abundance file, which is given in the input.in file with keyword nebComposition)
    cp /mocassin.X.XX/benchmarks/gas/HII40/abunHII40.in /mocassin.X.XX/input/abunHII40.in

2.1.3 Run a MOCASSIN pure photoionisation benchmark
The command you need to use to actually run MOCASSIN will depend on how your system is set up, and whether you are required to run your models through a queueing system, in which case you will probably need some sort of shell script to do that (ask your network administrator if unsure). If you want to run MOCASSIN outside any queueing system then you can try one of the following commands and hope that one works, if not ask you network administrator about how to interactively run MPI programs on your system.

for the SUN:
    mprun -np N ./mocassin (where N=number of processors you wish to use)

for the IBM
    mocassin -procs N (where N=number of processors you wish to use)

for the SGI, Beowulf and single Linux PCs using the standard MPI as above:
    mpirun -np N ./mocassin (where N=number of processors you wish to use)
2.1.4 Checking that everything is going OK

If the 'output' keyword is included in the input file (see 3.1), at the end of each iteration MOCASSIN will write out a number of output files (see 4.) in the output/ directory. Some of them are written out regardless of whether the 'output' keyword is included, these are grid files that are needed by the mocassinWarm driver (should you decide to stop the simulation and re-start it from the same iteration); see 3.1, 'writeGrid' keyword to see how to control the output grid files. Also MOCASSIN will keep you updated on what he is doing by outputting stuff to the screen (if that annoys you, just pipe it to a file e.g. mocassin -procs N > log; and then you can check the log at any time you wish). After each iteration it will tell you the percentage of grid-cells converged so far, as well as the number of 'dark cells' (i.e. those cells not reached by any of the photon packets) and also will provide a summary of eventual convergence problems generally this is the last bit of info written before the beginning of the new iteration, so, if you have decided to pipe the screen output into a file called log, then you can easily check the progress so far by typing tail -f log, or grep Summary log, or more simply, the convergence history is summarised in a file called 'output/summary.out.

For the benchmark cases you can start having an idea of whether MOCASSIN is working properly by waiting until it's about 60-70% converged and then checking the 'output/lineFlux.out' file. A description of this file is given in section 4. of this README file. Also check the 'output/temperature.out' and the 'output/ionratio.out' files (also described in section 4.).

2.1.5 Completing the benchmark runs

The benchmark input files included are set up to run a simulation of 13*13*13 grid cells (this is sufficient to reproduce the benchmark results; see Ercolano et al., 2003) and a given starting number of energy packets (e.g. this can be set by the keyword nPhotons 100000); one iteration of 100000 packet should run in a few minutes on most systems. The benchmarks input files make use of the Autopackets option (see 3.1), which which automatically increases the number of energy packets to be used as soon as a convergence plateau is reached.

This could also be done manually, by excluding the Autopackets option from the input file and starting off with about 1000000 packets. However 1000000 photons are not enough to reach full convergence in this case and you will have to stop the simulation when the percentage convergence reaches 30-40%, then edit the grid3.out file changing 1000000 to 3000000, and restart the simulation using the previously compiled mocassinWarm driver

    e.g. mprun -np N ./mocassinWarm

this will basically tell MOCASSIN to increase the number of energy packets to be used in the simulation from 1000000 to 3000000. It will, of course, now take longer to complete each iteration, but the convergence percentage should increase quite quickly. You can stop the simulation when you reach convergence > 90%. For simplicity, until you familiarise yourself with the program, it is advised that you keep the Autopackets option as set in the input files, in which case it will not be needed to manually stop and restart the simulation.

2.1.6 Comparing with the benchmark tables

The description of the output files is given in section 4. The output files contain all the info (and much more) you need to compare your simulation with the benchmark tables. Remember that MOCASSIN employs a statistical method, so do NOT expect to see exactly the same figures as those quote in the benchmark paper as some of the differences will be due to the statistical error and also to the convergence level/limit employed. Furthermore some of the atomic data has been updated since the publication of the paper.

2.2 Pure Dust Benchmarks

Version 2.0 was designed for the modelling of regions where dust grains and photoionised gas coexist in the same volume; however extreme care has been taken so that the code could also be run efficiently

when only one or the other process is included. In this section we will see briefly how to run dust-only models by attempting to reproduce pure dust 1D and 2D models as shown by Ercolano, Barlow and Storey (2005, MNRAS, in press).

2.2.1 1D Benchmarks
A set of spherically symmetric benchmark models and solutions are described by Ivezic et al. (1997, MNRAS, 291, 121). The input files used by MOCASSIN for some of these benchmarks are included in the directory /mocassin.X.XX/benchmarks/dust/1D.
Copy the input file of the benchmark you wish to run into the mocassin.X.XX/input subdirectory and rename it input.in or link the input.in file to the benchmark file. Then run the code (see Section 2.1.3) and finally compare the output files (SED.out dustGrid.out, see Section 4) with the results published by Ercolano, Barlow and Storey (2005, MNRAS, in press)

2.2.1 2D Benchmarks
A set of 2D disk benchmark models and solutions are described by Pascucci et al. (2004, A&A, 417, 793). Sample input files used by MOCASSIN for some of these benchmarks are included in the directory /mocassin.X.XX/benchmarks/dust/2D.
As for the 1D benchmarks you will have to copy the input files of the benchmark you wish to run into the mocassin.X.XX/input subdirectory and rename it input.in or link the input.in file to the benchmark file. Then run the code (see Section 2.1.3) and finally compare the output files (SED.out dustGrid.out, see Section 4) with the results published by Ercolano, Barlow and Storey (2005, MN-RAS, in press).

3. Running things other than benchmark cases

Once you have convinced yourself that MOCASSIN is behaving as it should, you can run your own simulations. If you are not confident that this is the case, please contact me before going any further.

The first step is to define your model. This is done via a set of keywords included in the input.in file. YOU SHOULD NEVER HAVE TO CHANGE ANYTHING IN THE SOURCE CODE (I hope). If you find you have to, please let me know. The rest of this section will list and describe all the keywords that can be included in the input file. The number of keywords you decide to use and the order you decide to list them in is entirely up to you. If you have missed out some required and fundamental keyword the simulation will stop and tell you what you have missed out. However, you should be aware that, for most of these keywords, default values are defined, so make sure that the default value is actually what you want before you decide to leave a given keyword out. Default values to each keyword are given below, enclosed by square brackets. The fundamental keywords, which must be specified in every input file are marked by an asterisk in the square brackets and have no default value.

3.1 List of keywords

autoPackets real1 real2 integer3: The number of energy packets to be used in the next iteration will automatically be increased by a factor of real2 whenever a convergence plateau (defined by real1) is reached, i.e. if the convergence level increase is less than the value specified by the user in real1. For example autoPackets 0.10 5. 1e8 will cause the total number of energy packets to increase by a factor of 5. whenever the convergence level increase from one iteration to the other is less than 10%. This only occurs when the convergence level is less than 85% and is the maximum number of packets defined by integer3 (1e8 in this example) has not been reached.
[.false.]

continuumCube real1 real2 : This keyword creates a 3D cube of the escaped continuum radiation in direction given by the inclination keyword and over all directions. If no inclination is specified in the input the continuum cube will include packets escaping in all directions. The continuum band wavelength limits are defined by real1 and real2 (give in μm). A negative value or the omission of this keyword will result in no continuum cube being written out.
[-1.]

contShape 'string' [optional: real]: The shape of the ionising continuum. The default value is 'blackbody', in which case the ionising stellar continuum is approximated by the Planck function for the stellar temperature defined by the keyword Stellar. If 'string' == 'powerlaw' then this must be followed by a real number indicating the index of the power law distribution such that $F_\nu \propto \nu^{-index}$. E.g. *contShape powerlaw 1.* will generate an input spectrum following a ν^{-1} distribution. Note that the final result will be normalised to the luminosity entered. The power law keyword is not compatiblewith the Lphot keyword, but only with the Lstar keyword. If a stellar atmosphere data file is to be used, the 'string' must specify the path of the external file containing the data. For example contShape NLTE140lg65 tells the program to look for the nLTe140lg65 data file in the current directory. The stellar atmosphere files must be in a format consisting of two columns: the first listing the wavelength points in units of [A] and the second containing the corresponding wavelength-dependent stellar Eddington fluxes in units of $[erg/cm^2/s/A/sr]$. A set of stellar atmosphere flux tables have been compiled by Dr T. Rauch in a MOCASSIN-friendly format and are available http://astro.uni-tuebingen.de/ rauch/ (but please manually remove the header from the flux tables.

[blackbody]

convLimit real : This is the convergence limit for the variation of a given parameter, in each grid cell from one Monte Carlo iteration to the next (e.g. 0.05 means changes of 5% maximum). In the case of a pure-gas (or gas+dust) simulation the criterion is based on the rate of change of neutral hydrogen. In the case of dust-only the criterion is based on the rate of change of the dust temperatures. Other convergence criteria can also be used, at the moment, this would require a simple editing of some source modules. If you would like to use a different convergence criterion please email me and I can do the editing for you.
[0.05]

debug : Logical switch to enable the debugging mode. When this keyword is included MOCASSIN will calculate a number of extra quantities (see Section 5.), which will, of course, slow the process down and also require more memory.
[.false.]

densityFile 'string' : The density structure of the nebula can be defined cell by cell by using an external density file. MOCASSIN knows that a density file is to be used when the densityFile 'string' is included in the input file, where 'string' contains the name and path of the file where the data is stored. This file must be consists of four, five or six columns, with the first three columns containing the x-, y-, and z- coordinates of the grid cell in [cm] and the fourth columns containing the value of the hydrogen density by number in $[cm^{-3}]$ at the particular grid cell. The x, y and z axis do not to be equally spaces, irregular grids are perfectly acceptable by MOCASSIN and also the extent of each axis can vary (as long as this is consistent with the values given in the nx, ny and nz fields). The fifth column is optional. If the multiChemistry keyword is specified the fifth column must contain an integer number in the range [1, Ncomponents] which indicates what component this cell belongs to (so that MOCASSIN can assign the chemical abundances for this component).
[]

densityLaw real... : This keyword is usually followed by a set of parameters which are to be fed into the density law routine, included in the grid_mod module. Any density law can be specified by editing the code in the setGrid subroutine. If the nebula is homogeneous, this keyword must be omitted and the Hdensity keyword included instead. Note that if neither of the two keywords is included, and an external density file is not specified with the densityFile keyword, the nebular density distribution is left undefined and the simulation halted with an error message being produced.
[0. 0. etc..]

dustFile string1 string2 : names dust data files -
string1 = grain species file
string2 = grain size info file
["none", "none"]

11

echo real1 real2 : If one needs to account for light travel time effects, for example if the source luminosity is changing, then this keyword will cause the results in SED.out to be integrated only over grid cells lying between ellipsoids corresponding to light travel times of real1 and real2. real1 and real2 are in light days, and real1 must be less than real2. The ellipsoids open out along the z-axis, and so one should also specify 'inclination 1 0.0 -1' in input.in when using this keyword. A review of the geometries of light echoes can be found in Sugerman, 2003, AJ, 126, 1939, and a case of MOCASSIN modelling including light travel time considerations is described by Wesson et al (arXiv:0907.0246). [.false.]

edges real1 real2 real3 : Defines the grid edges (only to be used with automatic grids). real1 real2 and real3 are respectively the x, y and z edges in cm
[-1., -1., -1. *must be given if automatic grids are used]

fillingFactor real1 : real1 can assume all values from 0. to 1. to defines the gas volume (and/or dust) filling factor
[1.]

getEquivalentTau : Logical switch to enable equivalent optical depth calculations (see Ercolano et al. 2007) - useful for diffuse/multiple and non-centrally located ionisation/ illumination sources
[.false.]

getTau : Logical switch to enable optical depth calculations and output - may be time consuming for large grids
[.false.]

Hdensity real : This keyword specifies a constant hydrogen density, by number, throughout the nebular region The command Hdensity 300 will e.g. set the hydrogen density, by number, to the constant value of 300 cm^{-3}.
[0.]

inclination integer real_1, real_2 ..... real_n1, real_n2 : This keyword controls the viewing angles at which the SED will be calculated as it will appear in the SED.out file (see Section 4.2). The number of viewing angles is given first (integer; in this version integer ¡= 2) and then the θ and ϕ inclination in radians. To turn off the ϕ dependence, real_2 must be set to to -1. . (θ =0. when the line of sight coincides with the z-axis)
[0]

inputNe : This indicates that the density distribution of the grid is in terms of electron densities instead of H density. This will cause the code to calculate at each iteration the values of H density from the local Ne values by taking into account the local ionisation structure.

[.false.]

isotropicScattering : This keyword activates the logical switch that turns off unisotropic scattering (implemented via the Heyney-Greenstein phase function - with ¡cos theta¿ calculated from the dielectric constants via MIE theory).
[.false.]

LPhot real : This is the number of hydrogen-ionizing photons emitted by the source per unit time, which is generally referred to as $Q(H^0)$, in the literature, with units of [E36 sec^{-1}]. If this is given then the stellar luminosity, LStar is automatically derived from it.
[* if LStar not given]

LStar real : This is the stellar luminosity in units of [E36 erg sec^{-1}]. If this is given as an input, then the number of hydrogen- ionizing photons, $Q(H^0)$, is automatically derived from it and from the input spectrum.
[* if LPhot not given]

maxIterateMC integer1 real1 : integer1 is the maximum number of Monte Carlo iterations to be performed in the simulation. real2 is the minimum convergence level to be achieved before the end of a simulation. The program will however stop after integer1 iteration even if real1% of convergence has yet to be reached.
[30 95.]

MdMg string1 real1/string2 : Dust to Gas ratio by mass. If string1 = 'constant' then it must be followed by real1, containing the value of MdMg to be applied homogeneously to all cells in the grid. If string1=file then it must be followed by string2, the name of the file defining the MdMg at each location. Note that MdMg, MdMh and Ndust are mutually exclusive. This file must consists of four columns, with the first three columns containing the x-, y-, and z- coordinates of the grid cell in [cm] and the fourth columns containing the dust to gas mass ratio at the particular grid cell. The x, y and z axis do not to be equally spaces, irregular grids are perfectly acceptable by MOCASSIN and also the extent of each axis can vary (as long as this is consistent with the values given in the nx, ny and nz fields).
[.false. 0./'none']

MdMh string1 real1/string2 : Dust to Hydrogen ratio by mass. If string1 = 'constant' then it must be followed by real1, containing the value of MdMh to be applied homogeneously to all cells in the grid. If string1=file then it must be followed by string2, the name of the file defining the MdMh at each location. Note that MdMg, MdMh and Ndust are mutually exclusive. This file must consists of four columns, with the first three columns containing the x-, y-, and z- coordinates of the grid cell in [cm] and the fourth columns containing the dust to hydrogen mass ratio at the particular grid cell. The x, y and z axis do not to be equally spaces, irregular grids are perfectly acceptable by MOCASSIN and also the extent of each axis can vary (as long as this is consistent with the values given in the nx,

ny and nz fields).
[.false. 0./'none']

multiChemistry integer1 string(1) string(2) ..... string(integer1) : This keyword must be used when a chemically inhomogeneous model is being performed. The integer1 value defined the number of components to be used. string(1), string(2) ... string(integer1) are the names of the files describing the abundances in each component. When the multiChemistry keyword is included the density distribution MUST be specified via the densityFile keyword. The fifth column of the density file must contain the index for the abundance file describing the chemical composition at each location. [.false.]

multiGrids integer1 string1 : This defines a multiple grid environment. The integer1 is the total number of grids to be used (mother + subgrids) and string1 is the name of the file containing the subgrid information. Please see the 'Running multiple spatial grids' section in this document. [.false. 'none']

multiPhotoSources string1 : This keyword is used to define multiple (or single) ionising sources , that can be placed at any locations. 'string1' is the name of the file containing the central star parameters. This file must contain in the first line the number of sources to be included and then each source should be specified on successive lines providing (in this order) T_{eff} [K], L_* [E36erg/s], ContShape, x-, y-, z-position. Note that the location of the source must be given normalised to the length of the relative axis. As example is included in the example/ directory.

nbins integer : The total number of points to be used in the frequency mesh. [600]

Ndust string1 real1/string2 : Number density $[cm^{-3}]$ of dust grains. If string1='constant' then it must be followed by real1, containing the value of Ndust to be applied homogeneously to all cells in the grid. If string1=file then it must be followed by string2, the name of the file defining Ndust at each location. Note that MdMg and Ndust are mutually exclusive. This file must consists of four columns, with the first three columns containing the x-, y-, and z- coordinates of the grid cell in [cm] and the fourth columns containing the number density of dust in $[cm^{-3}]$ at the particular grid cell. The x, y and z axis do not to be equally spaces, irregular grids are perfectly acceptable by MOCASSIN and also the extent of each axis can vary (as long as this is consistent with the values given in the nx, ny and nz fields). [.false. 0./'none']

nebComposition 'string' : This keyword specifies the path of the nebular composition data file. If the default 'solar' composition (defined in the file solar.dat) to be used, this keyword can be omitted. However the solar.dat file includes all elements with Z<30: this will result in a more memory expensive simulation. It is therefore advised to set to zero those elements which are not needed in the simulation. Otherwise the nebular composition can be specified in the user-defined 'string' file to be found in the

current directory. All composition input files must be in a format consisting of one column containing the abundances by number and relative to hydrogen for the first thirty elements in order of ascending atomic number. The abundances of elements which are not to be included in the simulations must be set to zero (this will automatically exclude them by flagging them out throughout the program). If the multiChemistry keyword is specified the nebComposition keyword must be omitted.
[* if gas is present and no multiChemistry]

NeStart real   : Initial guess for the nebular electron density.
[0.]

noPhotoelectric   : When this keyword is included in the input file all procedures associated with the calculation ofthe grain charges and photoelectric yield are switched off as well as gas-grain collision processes.
[$.false.$]

nPhotons integer   : This is the number of energy packets to be used in the Monte Carlo simulation and it has to be specified for each model.
[*]

nStages integer   : This is the number of ionisation stages to be used in the model.Max allowed is currently 10. If you have the atomic data necessary and would like to use more than 10 ionisation stages please contact me, or if you are confident you can edit the data/fileNames.dat and include the new data files to the poole.
[7]

nuMax real   : High limit of the frequency mesh in units of $[Ryd]$
[15.]

nuMin real   : Low limit of the frequency mesh in units of $[Ryd]$
[0.]

nx integer   : Number of axial points in the x-direction.
[30]

ny integer   : Number of axial points in the y-direction.
[30]

nz integer   : Number of axial points in the z-direction.
[30]

output   : when this keyword is included the output files will be written at the end of each iteration. If it is omitted no output will be created, however if the grid files are being created the output files can easily be recovered using the mocassinOutput driver.
[.*false.*]

planeIonization real1 real2   : this keyword is used when the ionisation source is from a plane and not from a point source. real1 must contain the value of the incident ionizing flux above nu = real2 [Ryd] (constant at each point on the plane) in units of $[phot/s/cm^2]$. If real2 =0. then the real1 is assumed to be the bolometric flux. When this keyword is specified the density distribution must be defined via the densityFile option. The ionizing plane is the x-z plane. the energy packets are reflected when they hit the y-z and the y-x planes and can escape from the x-z planes. (This can be easily changed to suit. please contact me be@star.ucl.ac.uk if your work requires it to be different)
[.*false.*]

quantumHeatGrain real1 real2   : This keyword activated the temperature spiking routines (quantum grain heating). It is only valid for simulations including dust. real1 is the limiting size of grain radii [um] that will be allowed to spike (i.e. grains with a ¡ real1 will spike). real2 is the minimum convergence level for the spiking to occur. Please read the notes on *Grain temperature spiking procedures* contained in this manual.
$[1.e − 3, 99.]$

quantumHeatGrainParameters real1 integer1 logical1   : This keyword provides controls over some internal parameters in the quantum grain heating procedures. They should not be modified unless you really know what you are doing. real1 is the max temperature to be considered in the grain temperature mesh of the spiking. integer1 is the number of temperature (and enthalpy) bins considered. logical1 switches on and off the writing out of a file containing all the probability functions for all the grains in every cell. The resulting file can be really gigantic and so this value should be set to .true. only for debugging purposes and used with care. Please read the notes on *Grain temperature spiking procedures* contained in this manual.
[700., 300, .*false.*]

Rin real   : Inner radius of the ionised region, in units of [cm].
[*]

Rout real   : Outer radius of the ionised region in units of [cm].
[0.]

recombinationLines : If this keyword is introduced, recombination line intensity of astrophysically abundant ions will be computed and appended to the lineFlux.out file
[.*false*.]

resLineTransfer real : real tells at what level of grid convergence the resonance line photons escape fractions should be calculated. This should be included when both dust and gas are present.
[101.]

slit real1 real2 : This keyword will cause the results in lineFlux.out, temperature.out and ionratio.out to be integrated over only those cell that fall under the projection of a slit aligned along the z-axis of the grid. The slit x and y dimensions (in [*cm*]) are defined by real1 and real2 respectively. If real1 and real2 have value 0. or if they are omitted, no slit is used and the results are integrated over the whole active volume.
[0., 0.]

symmetricXYZ : When the nebula to be modelled shows axial symmetry in the x- y- and z-directions, this keyword can be used to enable the symmetric grid procedures. This will result in the ionizing source being put in a corner of the grid, instead of being put in the centre, meaning that only one eighth of the nebula will have to be computed.
[.*false*.]

talk : This switch enables the verbose version of the program.
[.*false*.]

TeStart real : Initial guess for the nebular temperature.
[10000.]

traceHeating : Logical variable to switch on the thermal balance channel tracing. When this is included in the input file a file called *thermalBalance.out* will be written to the *output/* directory. Be aware that depending on the size of your grid this may be quite a large file and time-consuming in the I/O phase.
[.*false*.]

TStellar real : Temperature in [*K*] of the ionizing stellar source.
[*]

writeGrid real : real indicates the minimum grid convergence percentage after which the grid files will be written out.
[0.]

4. Input and Output Files

4.1 Input files

The source files are contained in a subdirectory called source/. MOCASSIN looks for the input.in file from a subdirectory called input/. The atomic data files should all be contained in a subdirectory named data/. Most of the atomic data files should not need to be changed at all. Unless you decide to update some of them, in which case (under the GPL agreement) you should also email me with the changes so that I can include them in the public version of the code. The dust optical data library and other dust related data files are contained in a subdirectory named dustData/.

The user's input files may be a combination of the following files, depending on the processes included in a given simulation.

input.in   This is the main input file where you can specify all the keywords to define your simulation. Some example input files are given for the Meudon/Lexington benchmark cases (see Section 2).

gas abundances file   This is the nebular abundances file which should have the name specified by the user in the nebComposition field of the input.in file. Some sample files are given for the Meudon/Lexington benchmarks.

gas density file   This the nebular density structure file which should have the name specified by the user in the densityFile field of the input.in file. The format of this file is given in Section 3 (see densityFile).

stellar atmosphere file   This is the stellar atmosphere file which should have the name specified by the user in the contShape field of the input.in file. The format of this file is given in Section 3 (see contShape).

dust number density file   This should contain the dust number density distribution across the grid. It's name and path should be specified in the input.in file by the keyword Ndust.

dust to hydrogen or dust to gas ratio   This contains the dust to hydrogen or dust to gas ratio distribution across the grid. It's name and path should be specified in the input.in file by the keyword MdMh or MdMg, respectively.

dust species and grain size distribution files   The names of these two files must be specified in the input.in file following the keyword dustFile. The dust species file should contain a first line specifying how many species are to be included, and then successive lines containing the names of the optical data (n,k or Qs) file and the relative abundance of the species.
e.g. for a pure silicate model (using the Draine and Lee 1984 data) :
1
'dustData/sil-dl.nk' 1.0
The grain size distribution file should contain a first line specifying how many grain sizes are to be included, the rest of the file should consist of three columns : index, radius (in um), weight. Grain size distribution files can be created using the makeGrainSizeDistribution.f90 program included in the accessories/ subdirectory.

e.g. for a single grain size
1 size
1 0.16 1.0

Input files for multigrid simulations are described in Section 5.2.

The plot.in file is used by the mocassinPlot driver in order to create 3D grids of line emission. This file must be place (or linked to) the input/ subdirectory. The plot.out and the grid4.out files are written out to output/ and can then be used to create emission line maps by integration along any given line of sight.
Monochromatic grids are created using the mono keyword, and individual lines using the line keyword. For example:

    mono
    line 2 4861. 4861.
    line 93 4686. 4686.
    line 1529 5007. 5007.
    line 1540 4363. 4363.
    line 2407 6733. 6733.
    line 2408 6718. 6718.
    line 929 6583. 6583.

The integer following the keyword line is the line index number as given in the lineFlux.out file. NOTE that the line indeces will be different for different simulations as they depend on which elements are included and on the number of ionisation stages accounted for. The 2nd and 3rd indeces contain the central wavelength of the line (these are redundant for monochromatic plots, however they must be included).

4.2 Output files

MOCASSIN produces several output files at various times during the simulation. This will be contained in a subdirectory named mocassin.X.XX/output/ The files ionratio.out, lineFlux.out, temperature.out, (tau.out), ionDen.out and SED.out are all produced by the output_mod module. The plot.out and grid4.out files are produced by the mocassinPlot file

ionratio.out   Ionratio.out contains the volume averaged ionic fractions. Different authors is the past have used slightly different definitions of this quantity in their models. Please refer to Ercolano et al. (2003) for further information on the description used by MOCASSIN.
The first two columns of the ionratio.out file give the atomic number of the element and its ionisation stage (1 for neutral, 2 for singly ionised etc.), and the third column gives the required quantity. If a multi-chemistry model is being run (see section 3.1), then the results will be given for each individual component.

lineFlux.out   The file lineFlux.out contains the volume integrated intensities of all the emission lines calculated by MOCASSIN. These are all given relative to H_β, which is in absolute units.
The first two columns give the element and ion number, these are followed by MOCASSIN codes for the levels of the transition; these are followed by the wavelength in [A]. The wavelength column is followed by the analytical and Monte Carlo line intensities relative to H_β, which is given in absolute units at the top of each region. Finally the last column gives the ratio of the two previous columns.

NOTE that the Monte Carlo line intensities are only calculated if the debug keyword is included in the input file. In normal mode only the intensities calculated using the formal solution (which are in general more accurate) will be available.

If a multi-chemistry model is being run (see section 3.1), then the results will be given for each individual component.

temperature.out   The file temperature.out contains the mean electronic temperatures weighted by the ionic species (see Ercolano et al., 2003, for definition). This file has the same structure as the ionratio.out file.

If a multi-chemistry model is being run (see section 3.1), then the results will be given for each individual component.

equivalentTau.out   Plese see Section 2.3 of Ercolano, Barlow and Sugerman (2007) for a description of this quantity. The first column contains Energy [Ryd] the second column contains wavelength [um] the third column contains equivalent tau [see paper] and the fourth column contains F^0_{lambda} [see paper]. Note that in the case of diffuse illumination this is the only meaningful quantity - see discussion in paper.

tauNu.out   The file tauNu.out contains the frequency dependent optical depth tau(nu) [which includes ALL opacity sources] at the edge of the grid in the three positive axial directions starting from the origin of the axes.

tau.out   THIS FILE IS NOT PRODUCED ANYMORE IN VERSIONS ¿= 2.02.49 PLEASE CONTACT BE IF YOU NEED IT. The file tau.out contains the run of the optical depth from the centre of the nebula to the outer edge along the three axial directions. The optical depths are calculated at the neutral hydrogen ionisation threshold, nu = 1.0 Ryd (13.6 eV), at the neutral Helium ionisation threshold, nu = 1.8 Ryd, and at the singly ionised Helium ionisation threshold, nu = 4.0 Ryd. The first column of the file gives the distance in [cm] from the centre of the nebula and the second column gives the optical depth from the centre to that point. This file is not always calculated correctly if the internal switches arenot set up properly. Please use equivalentTau. See description above.

ionDen.out   The file ionDen.out contains the ionic fractions at each grid cell. The first three columns give the x-, y- and z-axis indeces of the cell, the fourth and fifth columns give the atomic number and the ionisation stage of the element (as above, 1 for atom, 2 for singly ionised etc.) and, finally, the sixth column gives the corresponding ionic fraction.

SED.out   This file contains the emerging spectral energy distribution from the grid. As indicated in the files' header column 1 and column 2 contain the frequency [Ryd] and wavelength [um] grid, column 3 contains the direction averaged SED per unit direction (must multiply by Pi to obtain the total overall directions); the following columns contain the SED emerging in any given line of sight as requested by the user in the input.in file with the inclination keyword.

The grid files and photoSource.out   As we have already mentioned elsewhere, grid files are also written out after each iteration by routines contained in the grid_mod module. These are needed by the warm start driver (mocassinWarm) to re-initialise an interrupted simulation. These files are formatted such that they can be written out and read back in quickly and therefore they may not be very clear to the human eye. However, most of the information they contain is also given in a more intelligible form in the other output files listed above, dust temperatures are an exception as discussed in Section 5.5. Gas-only simulations will result in 4 grid files being written out: grid0.out, grid1.out, grid2.out and grid3.out.

The first line of the grid0.out file gives the number of grids included (mother+ subgrids), on the next line are the x-, y- and z-axes points in the mother grid, followed by the outer radius in [cm]. The next few lines list the x-axis points, then the y-axis points and, finally, the z-axis points. The rest of the file contains the convergence info for each grid cell. The active index of the cell in the first column, whether it has converged (1=yes,0=no), and whether it is a black cell (i.e. if could not be reached by any photons, 1=black, 0=normal). Cells that have a 0 index are inactive cells; cells with a negative index are cells that have been replaced by a subgrid, whose index is equal to the absolute value of the negative mother-cell index. In the Case of multigrid simulations, the file will loop around all subgrids included.

The grid1.out file contains the electron temperatures, electron densities and hydrogen densities for each grid cell in each grid. As for the grid0.out file, this information is given for each grid cell, with the last index varying the fastest (i.e. (1,1,1), (1,1,2), etc.).

The file grid2.out contains the ionic fractions at each grid cell for the ions included in the simulation. These are given in order of increasing atomic number and ionisation stage, with each element occupying one line. The grid cell indeces vary in the same fashion as in the grid1.out file.

The file grid3.out contains a list of the specified simulation parameters in a fixed order (the keywords are indicated on the right of each value). NOTE that it is not possible anymore to change nuStepSize from the input.in file. If you wish to change this parameter (you shouldn't need to), this is defined in the set_input_mod module.

Dust-only simulations only produce grid0.out, grid3.out and dustGrid.out

The file dustGrid.out contains the dust number density at each cell followed by the grain temperatures for each grain size of each grain species. For each cell Ndust is written on one line and this is followed by n_size+1 lines of n_species+1 columns containing the individual grain temperatures for each size and species, where n_size=number of size bins and n_species=number of grain species included. The average temperature for the grain mixture weighted by the size distribution and the species abundances is given at (0,0).

Dust+gas simulations will produce all the files above.

grid4.out is written out by mocassinPlot and it just contains the volume of each gridcell in $[e45cm^3]$, which is needed for visualisation purposes.

The ionising source(s) input parameters are written out to a file named photoSource.out.

plot.out   The MOCASSINPlot driver produces also an output file, containing the luminosities of each individual grid volume element in the required emission lines. This file, named plot.out, is written in a format which should be easily readable into a data visualisation software, such as IDL or PDL. A grid4.out is also written out, containing the volume of each gridcell in $[e45cm^3]$.

5. Miscellaneous notes on MOCASSIN

5.1 Analytical and Monte Carlo line fluxes

The total luminosity of the nebula in various emission lines longward of the Lyman limit can be obtained by using two methods. The first method, which is only available to Monte Carlo codes, consists of summing up the number of energy packets in the given line over all the grid cells. From this, the power emitted in that line can be readily obtained (see e.g. Ercolano et al., 2003). The second method consists of using the values of the local electron temperatures and ionic abundances given by the final model solution to obtain the line emissivities for each grid cell. The luminosity of the nebula in any given line can then be calculated easily by integrating the emissivity of the required line over the volume of the nebula.

A comparison of the results obtained using the two methods described above, provides an indication of the level of statistical accuracy achieved during the simulation, as the two methods will give consistent results only if enough energy packets are used in order to yield good statistics for every line. However, in general, the second method (formal solution) yields the most accurate results, particularly for weak lines, which only emit a few photons.

Calculation of the line emissivities using the first method, although straightforward, requires extra book-keeping which can be expensive for larger simulations. For this reason, this calculation will only be carried out when the keyword debug is included in the input file, otherwise the more speedy (and accurate) formal solution will only be employed.

5.2 Running multiple spatial grids

WARNING: This section is out of date – new help is in the process of being written as we speak – almost ;-) . Please email me if you want to use multiple grids and need some help be@star.ucl.ac.uk

From Version 2.00 onwards you can choose to run simulations including multiple spatial sampling. A typical example of when this would be need is for the modelling of knots embedded in a gaseous nebula. The density enhancement in the knot requires a finer spatial grid than one that may be sufficient to model the main nebula. In such cases the subgrid describing the knot can be modelled simultaneously and self consistently by MOCASSIN's multigrid routines. The radiation field will be safely transferred from mother to sub grids with (hopefully) no error being introduced by the process. The overheads involved only reflect the eventual introduction of extra grid cells.

A typical example of how such a model would be set up is briefly described below. All files mentioned should be included in the examples/multigrid subdirectory.

The example below shows the input.in file for a plane parallel model consisting of a cubic blob embedded into a cubic grid:

```
autoPackets 0.20 2. 1000000000
output
contShape blackbody
nebComposition 'examples/abunHII40.in'
maxIterateMC 30 95.
nPhotons 10000000
nx 10
ny 10
nz 10
planeIonization 3.0
Rin 0.
Rout 2.1e19
TStellar 45000.
```

```
writeGrid 5.
convLimit 0.03
densityFile 'examples/cubenew.dat'
multiGrids 2 'examples/subGrids.in'
```

The density distribution of the mother grid (containing the main nebula) is provided provided by
an external file using the densityFile keyword as described in Section 3.1; e.g.

```
densityFile 'examples/cube.dat'
```
MOCASSIN knows that it will have to deal with more than one grid (in this case 2) thanks to the line
```
multiGrids 2 'examples/subGrids.in'
```
the integer after the multiGrid keyword is the total number of grids including the mother grid; the
file 'example/subGrids.in' contains all the information regarding the position of the subgrids and their
specification. This file is created automatically by using the makeSubGrids.f90 program included in
the 'accessories/' directory.
The makeSubGrids.f90 program reads the makeSubGrids.in input file, that has the following form

```
1
cube.dat 10 10 10 1
cubenew.dat
1 9 9 9 knot_den.in
7.e18 14.e18 7.e18 14.e18 7.e18 14.e18
100. -1
```

line 1: number of subgrids to be included
line 2: name of mother grid file; its x,y,z dimensions; multiChemistry? (1=no; >1=yes)
line 3: name of the modified mother grid file (to be created by makeSubGrids.f90)
line 4: knot index; its x,y,z extent; name of knot file (to be created by makeSubGrids.f90)
line 5: xmin,xmax,ymin,ymax,zmin,zmax in cm, defining the region occupied by knot 1
line 6: H number density in $[cm^{-3}]$ of knot 1, abundance index (negative if no multiChemistry)

The maximum number of subgrids that may be included in a simulation is controlled by the pa-
rameter maxGrids in the source/constants_mod.f90 module; in the example above only one subgrid is
included, additional subgrids could be included by simply repeating lines 4,5 and 6 for each subgrid.

The knot_den.in file above has the following format

```
0.0000000E+00 0.0000000E+00 0.0000000E+00 100.0000
0.0000000E+00 0.0000000E+00 0.1250000 100.0000
0.0000000E+00 0.0000000E+00 0.2500000 100.0000
0.0000000E+00 0.0000000E+00 0.3750000 100.0000
0.0000000E+00 0.0000000E+00 0.5000000 100.0000
0.0000000E+00 0.0000000E+00 0.6250000 100.0000
0.0000000E+00 0.0000000E+00 0.7500000 100.0000
0.0000000E+00 0.0000000E+00 0.8750000 100.0000
0.0000000E+00 0.0000000E+00 1.000000 100.0000
0.0000000E+00 0.1250000 0.0000000E+00 100.0000
0.0000000E+00 0.1250000 0.1250000 100.0000
etc ........
```

Columns 1, 2 and 3 contain the x, y, and z positions in the subgrid, normalised to 1. The fourth column contains the h number density. (when running a multiChemistry model a fifth column would appear; containing the abundance file index).

It is worth noticing that since the x y and z positions in the knot_den.in files are given normalised to unity, we can include knots with the same geometry and density at several grid locations without the] need to create a density file for each knot. this can be done by specifying the same filename for the knots in the makeSubgrids.in file.

The makeSubgrids.f90 file will also process the cube.dat file (which is the original mother grid density file provided by the user) turning off all the cells that are to be replaced by subgrids. The new file will be called cubenew.dat in the above example and this is the file that must be specified in the input.in file of mocassin.

The makeSubgrids.f90 file can only create subgrids of homogeneous density and chemical indeces. It should be very easy for the user to customise the grids obtained to included their 'favourite' density/chemistry distribution.

The output from multigrid models will be summarised in the usual files and an overall, as well as grid by grid result will be provided.

Christophe Morisset of UNAM is currently in the process of creating an IDL visualisation tool that is able to deal with multiGrids and reconstruction for a quick and efficient analysis of the results. It is intended to distribute the new tool as soon as it becomes available.

5.3 Running models that include dust and gas.

The gas and dust radiative transfer routines in MOCASSIN are now fully integrated. It is therefore now possible to run MOCASSIN in its original gas-only mode, as well as dust-only and of course dust+gas.

The basics on how to run the code for dust-only or gas-only cases have already been given in Section 2 above; here we will concentrate on an example where both dust and gas are present.

Below is the input.in file used for the dust and gas model of NGC 3918 as described by Ercolano, Barlow and Storey (2005, MNRAS, in press).

```
autoPackets 20. 2. 500000000
densityFile 'ngc3918/ngc3918den.dat'
symmetricXYZ
multiChemistry 2 'ngc3918/ngc3918.dat' 'ngc3918/ngc3918.dat'
contShape 'ngc3918/nLTe140lg65'
maxIterateMC 30 95.
nPhotons 1000000
nx 23
ny 23
nz 23
nbins 700
LStar 27.64
nuMax 23.7
nuMin 3.1e-4
Rin 0.
Rout 3.27142e17
TStellar 140000.
MdMh constant 0.0011
dustFile 'ngc3918/primary_grainspecies.dat' 'ngc3918/primary_grainsizes.dat'
```

```
writeGrid 50.
convLimit 0.03
resLinesTransfer 90.
```

In summary, it should be clear from the example above that the only keywords that differentiate the file above from the input of a gas-only model are : MdMh (dust to hydrogen mass ratio), dustFile (files containing the grain size distribution and species information) and resLineTransfer (which tells at what level of grid convergence the resonance line photons escape fractions should be calculated). Note that the first two keywords that define how much, what type and what size grains are to be used is also needed to run dust-only models (although MdMh could be substituted by Ndust or by MdMg as described in Section 3). resLinesTransfer is only needed if there is gas also in the simulation (which would then be emitting resonance lines capable of heating the grains).

5.4 The accessories/ subdirectory

A number of useful (or not) FORTRAN and IDL programs are included in the accessories/ subdirectory. Some words of warning: the programs are very basic and poorly commented, as they were developed for personal use. Anyone is welcome to use them at their own risk!! The IDL programs, in particular, are specific to the simulation they were developed for and some editing will be necessary to customise them to the user's needs.

5.5 Plotting dust temperatures

Depending on the complexity of the dust model employed in a given model and the geometrical complexity of the grid, it may be quite challenging to explore the dust temperature distribution calculated by mocassin and written out to dustGrid.out (see Section 4.2).

Sometimes the only way is to plot out the results or create 3D maps using a visualisation program such as IDL, PDL etc..

The accessories/ subroutine contains an example (dustTemperatures2.pro) of how such a grid may visualised using IDL. This was written for a grid containing 2 grain species and 20 grain sizes. The simulation was a multiChemistry dust and gas one, so the results are also split by sector.

5.6 Grain temperature spiking routines

The grain temperature spiking routines included in MOCASSIN are based on the Guhathakurta & Draine (1989 ApJ 345, 230) method. This is a very powerful method and allows the time-efficient computation of the time-dependent grain temperatures due to quantum heating. For the limitations of the method also see Siebenmorgen et al. (1992 A&A 266, 501). The temperature spikes only affect the output SED from dust grains, it is therefore advisable not to include these time consuming procedures until the model has almost converged in the case of gas+grain simulations (keep a high value of real2 - see keyword description). In the gas of dust only models, it is worth to have the procedures working right from the start since the convergence criterion is then based on dust temperatures and therefore one must take this effect into account right from the start in order to avoid copnvergence fluctuations. It is in general not worth running the quantum heating routines on large grains that are unlikely to spike. For a discussion of the general cases when quantum heating routines must be considered pls see Siebenmorgen et al. (1992 A&A 266, 501).

At present the grain temperature spiking routines are only implemented for carbonaceous or silicate grains. _mocassin_ will expect to be told what type of grain he is dealing with when calculating the spiking. This is done by adding a -capital- 'S' or 'C' at the beginnig of the species label in the optical constants file. e.g. for the Draine & Laor (1993) silicates data (dustData/sil-dlaor.nk):

nk

Ssil_dl 1400. 3.3 0.588 20.077

0.10000E-02 0.99956E+00 0.97380E-04

0.10120E-02 0.99955E+00 0.10160E-03

0.10230E-02 0.99954E+00 0.10610E-03

0.10350E-02 0.99953E+00 0.11060E-03

0.10470E-02 0.99952E+00 0.11520E-03

0.10590E-02 0.99951E+00 0.12000E-03

....

Please email (be@star.ucl.ac.uk) me if you would like to include T-spiking effects for other species.

6. Limitations and future development

MOCASSIN's principal limitation is imposed by the computer power available. The great volume of data which has to be handled in a three-dimensional simulation, implies the need for a system with multi-processing capabilities in order to accelerate the computational time. However, the fast development of Beowulf Linux clusters is making parallel computing more affordable, and this is also a reason why the MPI formalism was chosen (as opposed to openMP), as this allows information to be passed from one processor to another and, therefore, it does not necessarily require a system with shared memory facilities. Such systems, which include the Silicon Graphics Origin 2000 machine used for this work, are generally much more expensive than Beowulf clusters.

Pure-dust models are less computationally expensive than gas or dust and gas ones and reasonably sizes grids can be feasibly run on single processor machines.

MOCASSIN was designed for the modelling of the photoionised region of planetary nebulae and H II regions and it does not, at present, include the high energy physical processes which are needed, for example, for the modelling of AGNs. However the inclusion of processes such as inner shell photoionisation and Compton heating is straightforward and this is intended to be one of the developments of the near future. Moreover, current efforts are geared toward the inclusion of a self-consistent treatment for PDR processes, this requires the incorporation of a chemical network.